

# DE300 HW1 - Nicky Williams

```
In [2]: import pandas as pd
inventory = pd.read_csv('/content/T_F41SCHEDULE_B43.csv')
# inventory.head()
```

```
<ipython-input-2-1067a66527ff>:2: DtypeWarning: Columns (11) have mixed type
s. Specify dtype option on import or set low_memory=False.
inventory = pd.read_csv('/content/T_F41SCHEDULE_B43.csv')
```

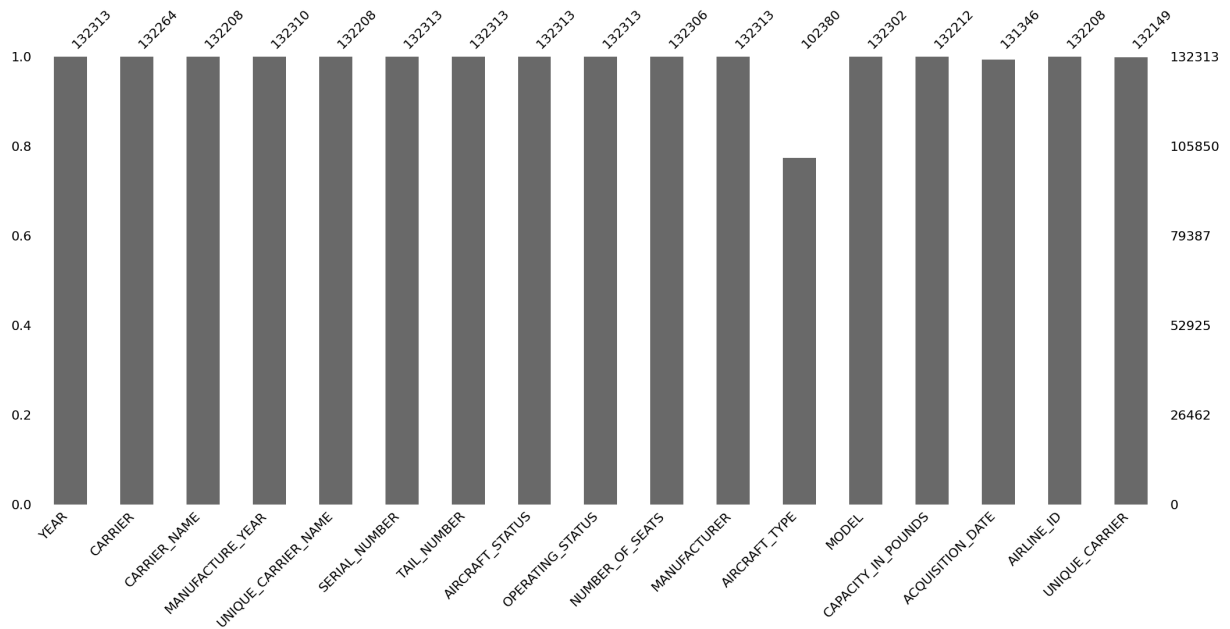
1. Investigate the missing data in this dataset. Specifically, for each of the following variables that have missing data, decide if any imputation is possible. Give your reasoning and code if you decide to impute missing values.
  - Columns for investigation: CARRIER, CARRIER\_NAME, MANUFACTURE\_YEAR, NUMBER\_OF\_SEATS, CAPACITY\_IN\_POUNDS, and AIRLINE\_ID.
  - For example, watch out for "North American Airlines" aircrafts. Are the CARRIER/UNIQUE\_CARRIER column really missing?

```
In [12]: # Inspecting missing values in given columns
cols_to_check = ['CARRIER', 'CARRIER_NAME', 'MANUFACTURE_YEAR',
                 'NUMBER_OF_SEATS', 'CAPACITY_IN_POUNDS', 'AIRLINE_ID']
missing_summary = inventory[cols_to_check].isnull().sum()
print("Missing values in each column:")
print(missing_summary)
```

```
Missing values in each column:
CARRIER          0
CARRIER_NAME     0
MANUFACTURE_YEAR  0
NUMBER_OF_SEATS   0
CAPACITY_IN_POUNDS 0
AIRLINE_ID        0
dtype: int64
```

```
In [4]: import missingno as msno
msno.bar(inventory)
```

```
Out[4]: <Axes: >
```



```
In [5]: # Note that the "NA" in the North American Airlines aircrafts should not be
# as missing columns but rather as a just poor choice of a shorthand notation

# Impute carrier using carrier_name because they have a one-to-one relations
carrier_mapping = inventory[['CARRIER', 'CARRIER_NAME']].dropna().drop_duplicates()
name_to_carrier = dict(zip(carrier_mapping['CARRIER_NAME'], carrier_mapping['CARRIER']))
inventory['CARRIER'] = inventory.apply(
    lambda row: name_to_carrier.get(row['CARRIER_NAME'], row['CARRIER']) if
    axis=1
)

# Impute carrier_name using carrier (same logic as above)
code_to_name = dict(zip(carrier_mapping['CARRIER'], carrier_mapping['CARRIER_NAME']))
inventory['CARRIER_NAME'] = inventory.apply(
    lambda row: code_to_name.get(row['CARRIER'], row['CARRIER_NAME']) if pd.
    axis=1
)

# Impute manufacture_year using model because the manuf. year is usually the
# across the same aircraft models minus a few outliers (hence using median)
# Fill any remaining missing values with the median value of the column
inventory['MANUFACTURE_YEAR'] = inventory.groupby('MODEL')['MANUFACTURE_YEAR'].apply(
    lambda x: x.fillna(x.median())
)
inventory['MANUFACTURE_YEAR'] = inventory['MANUFACTURE_YEAR'].fillna(inventory['MANUFACTURE_YEAR'].median())

# Impute number_of_seats using model (same logic as above)
inventory['NUMBER_OF_SEATS'] = inventory.groupby('MODEL')['NUMBER_OF_SEATS'].apply(
    lambda x: x.fillna(x.median())
)
inventory['NUMBER_OF_SEATS'] = inventory['NUMBER_OF_SEATS'].fillna(inventory['NUMBER_OF_SEATS'].median())

# Impute capacity_in_pounds using model (same logic as above)
inventory['CAPACITY_IN_POUNDS'] = inventory.groupby('MODEL')['CAPACITY_IN_POUNDS'].apply(
    lambda x: x.fillna(x.median())
)
```

```

)
inventory['CAPACITY_IN_POUNDS'] = inventory['CAPACITY_IN_POUNDS'].fillna(inv

# Impute airline_id using carrier because they have a one-to-one relationshi
airline_map = inventory[['CARRIER', 'AIRLINE_ID']].dropna().drop_duplicates(
carrier_to_id = dict(zip(airline_map['CARRIER'], airline_map['AIRLINE_ID']))
inventory['AIRLINE_ID'] = inventory.apply(
    lambda row: carrier_to_id.get(row['CARRIER'], row['AIRLINE_ID']) if pd.i
    axis=1
)

# Check
print("Final missing values:")
print(inventory[['CARRIER', 'CARRIER_NAME', 'MANUFACTURE_YEAR',
                  'NUMBER_OF_SEATS', 'CAPACITY_IN_POUNDS', 'AIRLINE_ID']].isr

```

Final missing values:

CARRIER	0
CARRIER_NAME	0
MANUFACTURE_YEAR	0
NUMBER_OF_SEATS	0
CAPACITY_IN_POUNDS	0
AIRLINE_ID	0

dtype: int64

2. Inspect the columns MANUFACTURER, MODEL, AIRCRAFT\_STATUS, and OPERATING\_STATUS. Decide, for each column, if transformation or standardization of data are required. Give your reasoning and code if you decide to transform the data. **Hints:**

- For very messy data like manufacturer/model names, give your best attempt. It is okay to not catch them all.
- Use value\_counts() to identify "big wins".
- Break down into multiple steps, instead of having one line of code to do them all.

The manufacturer column needs standardization because it has several variations for the same company due to inconsistent capitalization, spacing, and naming conventions (BOEING/THEBOEINGCO/BoeingCo/THEBOEINGCOMPANY, etc).

The model column *may* need standardization depending on how exact the models are or if letters/words can be included in the model id.

The aircraft\_status and operating\_status columns have discrepancies in capitalization and need standardization for that.

```

In [6]: for col in ['MANUFACTURER', 'MODEL', 'AIRCRAFT_STATUS', 'OPERATING_STATUS']:
        print(f"\nUnique values in {col}:")
        print(inventory[col].value_counts(dropna=False).head(20))

```

Unique values in MANUFACTURER:

MANUFACTURER

BOEING

15922

Embraer

11508

THEBOEINGCO

9223

Bombardier

8871

Boeing

8392

BoeingCo

7446

AIRBUS

7179

AirbusIndustries

6967

BOEINGCOMPANY

6767

Airbus

5289

CESSNA

4157

EMBRAER

3287

MCDONNELL-DOUGLAS

3160

BOMBARDIER

2821

BOEING

2811

MCDONNELLDUGLAS

2781

CANADAIR

2734

AirbusIndustrie

2666

THEBOEINGCOMPANY

2142

TheBoeingCompany

1833

Name: count, dtype: int64

Unique values in MODEL:

MODEL

EMB-145 2614

B-737-7H4 2470

B737-823 2370

A320-232 2333

A321-231 2259

737-700PASSENGERONLY 2027

C-208B 1872

B757-2 1775

CRJ-2/4 1761

B737-800PAX 1621

MD-80	1610
ERJ-170-200LR	1379
757-200	1345
CRJ200-2B19	1342
A319	1267
B-737-8H4	1256
CRJ-200	1148
ERJ-175	1132
SUPER80PASSENGER	1108
A320-1/2	1107

Name: count, dtype: int64

Unique values in AIRCRAFT\_STATUS:

AIRCRAFT_STATUS	
0	79487
b	30852
B	12699
a	7804
A	1330
L	122
o	19

Name: count, dtype: int64

Unique values in OPERATING\_STATUS:

OPERATING_STATUS	
Y	126577
N	5664
y	71
	1

Name: count, dtype: int64

```
In [7]: import numpy as np

# 1. Manufacturer col

# Standardizing whitespace and casing
inventory['MANUFACTURER'] = inventory['MANUFACTURER'].str.strip().str.upper()

# Mapping common variations of "Boeing" to standardized names
manufacturer_map = {
    "BOEING": "BOEING",
    "THEBOEINGCO": "BOEING",
    "BOEINGCO": "BOEING",
    "BOEINGCOMPANY": "BOEING",
    "THEBOEINGCOMPANY": "BOEING",
    "THEBOEINGCOMPANY": "BOEING",
    "THEBOEING": "BOEING",
    "THEBOEING CO": "BOEING",
    "THEBOEINGCOMP": "BOEING",
    "THEBOEING COMPANY": "BOEING",
    "THEBOEINGCOMPANY": "BOEING",
    "THEBOEINGCO": "BOEING",
    "THEBOEINGCOMP": "BOEING",
    "BOEING COMPANY": "BOEING",
    "AIRBUS": "AIRBUS",
    "AIRBUSINDUSTRIES": "AIRBUS",
```

```

    "AIRBUSINDUSTRIE": "AIRBUS",
    "EMBRAER": "EMBRAER",
    "BOMBARDIER": "BOMBARDIER",
    "CESSNA": "CESSNA",
    "MCDONNELL-DOUGLAS": "MCDONNELL DOUGLAS",
    "MCDONNELLDUGLAS": "MCDONNELL DOUGLAS"
}

inventory['MANUFACTURER'] = inventory['MANUFACTURER'].replace(manufacturer_map)

# 2. Model col

# Standardizing spacing and casing
inventory['MODEL'] = inventory['MODEL'].str.strip().str.upper()

# Some obvious variants
model_map = {
    "737-700PASSENGERONLY": "737-700",
    "SUPER80PASSENGER": "MD-80",
    "B737-800PAX": "B737-800",
    "CRJ200-2B19": "CRJ-200",
    "CRJ-2/4": "CRJ-200",
    "A320-1/2": "A320"
}

inventory['MODEL'] = inventory['MODEL'].replace(model_map)

# Aircraft_status col
# Standardize casing
inventory['AIRCRAFT_STATUS'] = inventory['AIRCRAFT_STATUS'].str.upper()

# Operating_status col
# Standardize casing
inventory['OPERATING_STATUS'] = inventory['OPERATING_STATUS'].str.upper()
# Treat empty strings or whitespace-only as missing
inventory['OPERATING_STATUS'] = inventory['OPERATING_STATUS'].replace(r'^\s*', '')

```

```

In [8]: # Checking
print("Unique cleaned values:")
print("\nMANUFACTURER:", inventory['MANUFACTURER'].value_counts().head())
print("\nMODEL:", inventory['MODEL'].value_counts().head())
print("\nAIRCRAFT_STATUS:", inventory['AIRCRAFT_STATUS'].value_counts())
print("\nOPERATING_STATUS:", inventory['OPERATING_STATUS'].value_counts(dropna=False))

```

Unique cleaned values:

```
MANUFACTURER: MANUFACTURER
BOEING          54933
AIRBUS          23159
EMBRAER         15554
BOMBARDIER      11834
MCDONNELL DOUGLAS 8465
Name: count, dtype: int64
```

```
MODEL: MODEL
CRJ-200        4251
EMB-145        2976
MD-80          2718
B-737-7H4      2470
B737-823       2370
Name: count, dtype: int64
```

```
AIRCRAFT_STATUS: AIRCRAFT_STATUS
0      79506
B      43551
A       9134
L        122
Name: count, dtype: int64
```

```
OPERATING_STATUS: OPERATING_STATUS
Y      126648
N       5664
NaN         1
Name: count, dtype: int64
```

3. Remove data rows that still have missing values. Report the amount of remaining data you obtained.

```
In [9]: # Dropping rows w/ any remaining missing values and printting the result
inventory_cleaned = inventory.dropna()
print(f"Number of rows after dropping missing values: {inventory_cleaned.sha
```

Number of rows after dropping missing values: 101275

#### 4. Transformation and derivative variables

- For the columns NUMBER\_OF\_SEATS and CAPACITY\_IN\_POUNDS, check the skewness in the variable and plot a histogram for each variable.
- The Box-Cox transformation (`scipy.stats.boxcox`) is one possible way to transform variables into a "more-normal-like" variable. Apply the Box-Cox transformation for these two columns and save them as new columns, i.e. XXXXXXXXXX\_BOXCOX.
- Plot a histogram for each transformed variable.
- Describe what you observe before and after transformation.

```
In [10]: import matplotlib.pyplot as plt
from scipy.stats import skew

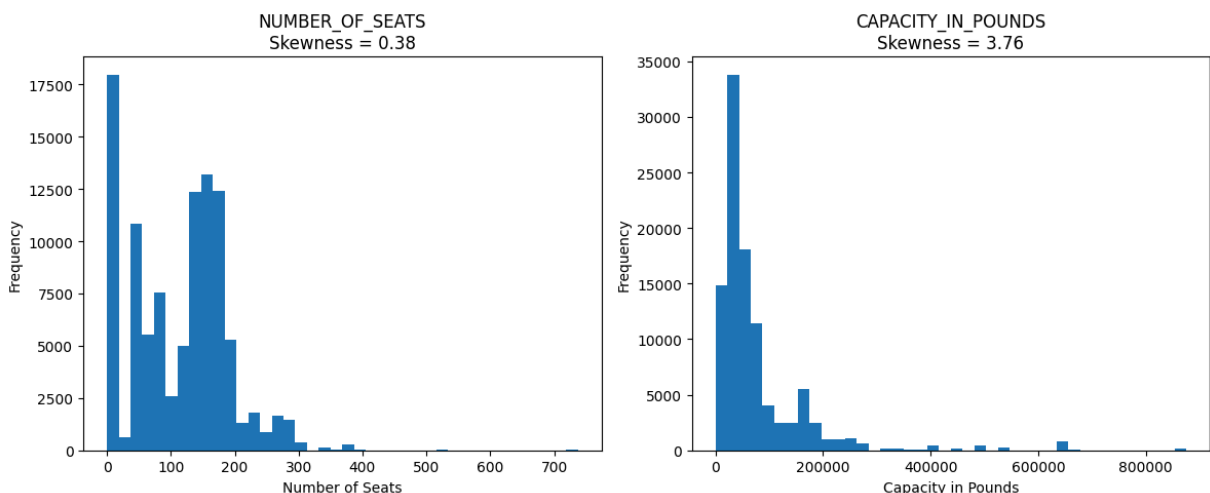
# Calculating skewness
seats_skew = skew(inventory_cleaned['NUMBER_OF_SEATS'])
capacity_skew = skew(inventory_cleaned['CAPACITY_IN_POUNDS'])

# Plotting histograms
plt.figure(figsize=(12, 5))

# Number_of_seats
plt.subplot(1, 2, 1)
plt.hist(inventory_cleaned['NUMBER_OF_SEATS'], bins=40)
plt.title(f'NUMBER_OF_SEATS\nSkewness = {seats_skew:.2f}')
plt.xlabel('Number of Seats')
plt.ylabel('Frequency')

# Capacity_in_pounds
plt.subplot(1, 2, 2)
plt.hist(inventory_cleaned['CAPACITY_IN_POUNDS'], bins=40)
plt.title(f'CAPACITY_IN_POUNDS\nSkewness = {capacity_skew:.2f}')
plt.xlabel('Capacity in Pounds')
plt.ylabel('Frequency')

plt.tight_layout()
plt.show()
```



```
In [31]: from scipy.stats import boxcox

inventory_cleaned = inventory.dropna().copy()

# Using only rows that are positive since Box-Cox requires > 0
positive_seats = inventory_cleaned['NUMBER_OF_SEATS'] > 0
positive_capacity = inventory_cleaned['CAPACITY_IN_POUNDS'] > 0
valid_rows = positive_seats & positive_capacity

# Box-Cox
seats_boxcox, seats_lambda = boxcox(inventory_cleaned.loc[valid_rows, 'NUMBER_OF_SEATS'])
capacity_boxcox, capacity_lambda = boxcox(inventory_cleaned.loc[valid_rows, 'CAPACITY_IN_POUNDS'])
```



```
# Storing into new cols
inventory_cleaned.loc[valid_rows, 'NUMBER_OF_SEATS_BOXCOX'] = seats_boxcox
inventory_cleaned.loc[valid_rows, 'CAPACITY_IN_POUNDS_BOXCOX'] = capacity_boxcox

# Printing lambdas
print(f"Box-Cox lambda for NUMBER_OF_SEATS: {seats_lambda:.4f}")
print(f"Box-Cox lambda for CAPACITY_IN_POUNDS: {capacity_lambda:.4f}")
```

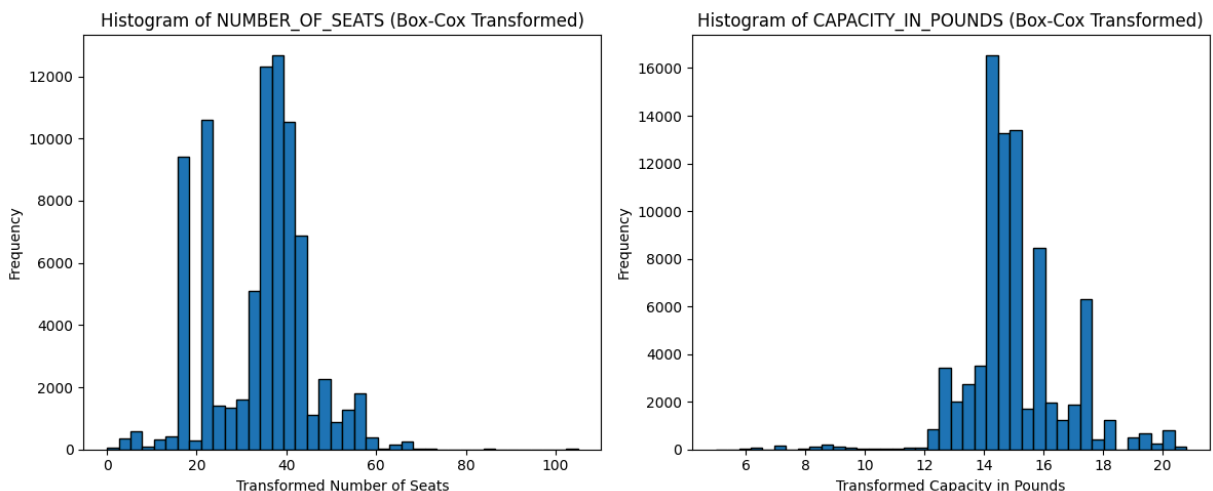
Box-Cox lambda for NUMBER\_OF\_SEATS: 0.6392  
Box-Cox lambda for CAPACITY\_IN\_POUNDS: 0.0575

In [33]: `import matplotlib.pyplot as plt`

```
# Creating histograms
plt.figure(figsize=(12, 5))

# Number_of_seats_boxcox
plt.subplot(1, 2, 1)
plt.hist(inventory_cleaned['NUMBER_OF_SEATS_BOXCOX'].dropna(), bins=40, edgecolor='black')
plt.title('Histogram of NUMBER_OF_SEATS (Box-Cox Transformed)')
plt.xlabel('Transformed Number of Seats')
plt.ylabel('Frequency')

# Capacity_in_pounds_boxcox
plt.subplot(1, 2, 2)
plt.hist(inventory_cleaned['CAPACITY_IN_POUNDS_BOXCOX'].dropna(), bins=40, edgecolor='black')
plt.title('Histogram of CAPACITY_IN_POUNDS (Box-Cox Transformed)')
plt.xlabel('Transformed Capacity in Pounds')
plt.ylabel('Frequency')
plt.tight_layout()
plt.show()
```



Before the Box-Cox transformation, both columns showed strong right-skewed distributions.

Number\_of\_seats was slightly more moderate, with most aircraft clustering around 50-200 seats and a few outliers with a large number of seats stretching the right tail of the distribution.

Capacity\_in\_pounds was very right-skewed, with an even smaller number of outliers (likely cargo planes).

After the Box-Cox transformation, both variables became more symmetric and closer to a normal distribution.

## 5. Feature engineering

- Create a new column SIZE by the quartiles of NUMBER\_OF\_SEATS
  - below 25% percentile: SMALL
  - 25% - 50% percentile: MEDIUM
  - 50% - 75% percentile: LARGE
  - above 75% percentile: XLARGE
- For each size group, provide and plot the proportions of aircrafts that are operating versus not (OPERATING\_STATUS).
- For each size group, provide and plot the proportions of aircrafts belonging to each aircraft status group (AIRCRAFT\_STATUS).
- Provide a written summary of your findings.

```
In [34]: # Quantiles
q1 = inventory_cleaned['NUMBER_OF_SEATS'].quantile(0.25)
q2 = inventory_cleaned['NUMBER_OF_SEATS'].quantile(0.50)
q3 = inventory_cleaned['NUMBER_OF_SEATS'].quantile(0.75)

def assign_size(seats):
    if seats < q1:
        return 'SMALL'
    elif seats < q2:
        return 'MEDIUM'
    elif seats < q3:
        return 'LARGE'
    else:
        return 'XLARGE'

# Creatting new col
inventory_cleaned['SIZE'] = inventory_cleaned['NUMBER_OF_SEATS'].apply(assign_size)
inventory_cleaned[['NUMBER_OF_SEATS', 'SIZE']].head()
```

```
Out[34]:
```

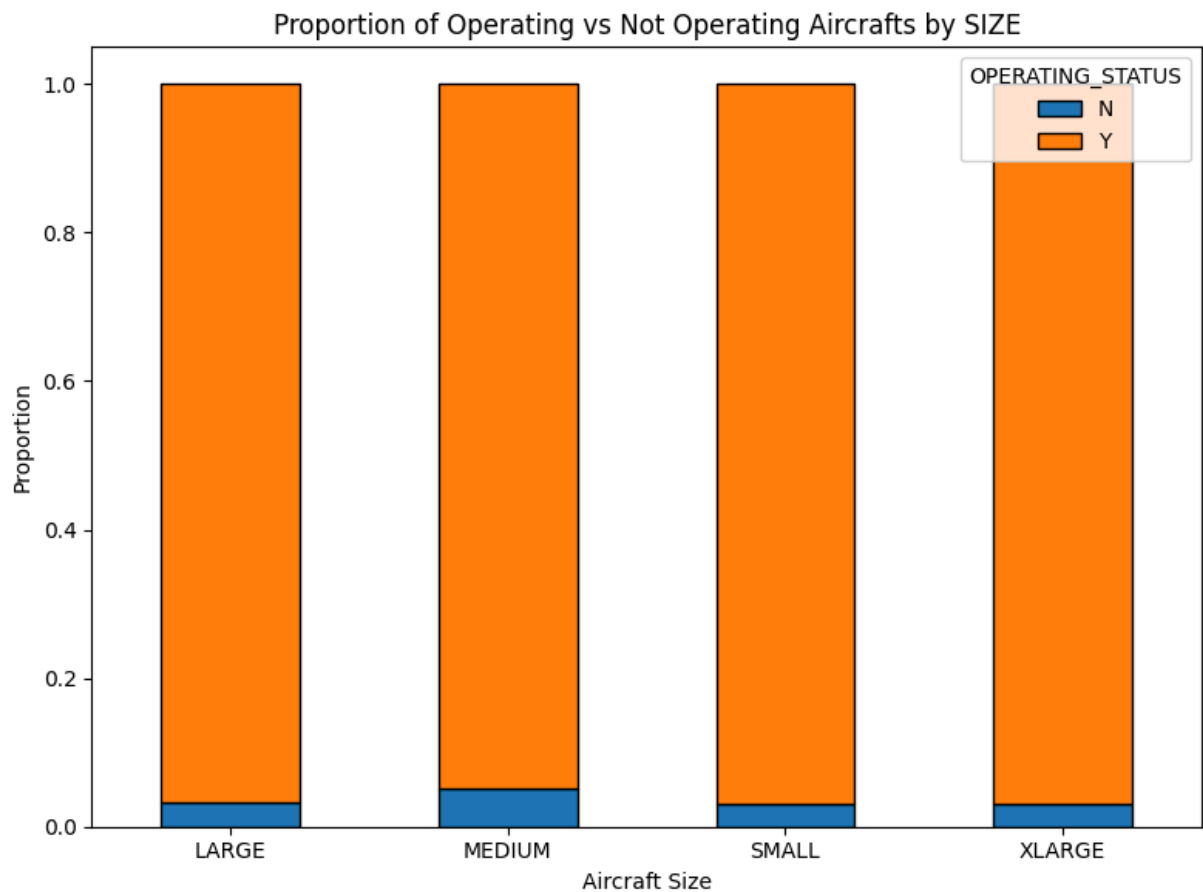
	NUMBER_OF_SEATS	SIZE
29239	92.0	MEDIUM
29240	86.0	MEDIUM
29241	86.0	MEDIUM
29242	136.0	MEDIUM
29243	19.0	SMALL

```
In [35]: import matplotlib.pyplot as plt

# Grouping by size and operating_status
operating_summary = (
    inventory_cleaned.groupby(['SIZE', 'OPERATING_STATUS'])
    .size()
    .unstack(fill_value=0)
)

# Converting to proportions
operating_proportions = operating_summary.div(operating_summary.sum(axis=1),

# Plotting
operating_proportions.plot(kind='bar', stacked=True, figsize=(8, 6), edgecolor=
plt.title('Proportion of Operating vs Not Operating Aircrafts by SIZE')
plt.xlabel('Aircraft Size')
plt.ylabel('Proportion')
plt.legend(title='OPERATING_STATUS')
plt.xticks(rotation=0)
plt.tight_layout()
plt.show()
print(operating_proportions)
```



OPERATING_STATUS	N	Y
SIZE		
LARGE	0.033143	0.966857
MEDIUM	0.052408	0.947592
SMALL	0.031002	0.968998
XLARGE	0.030592	0.969408

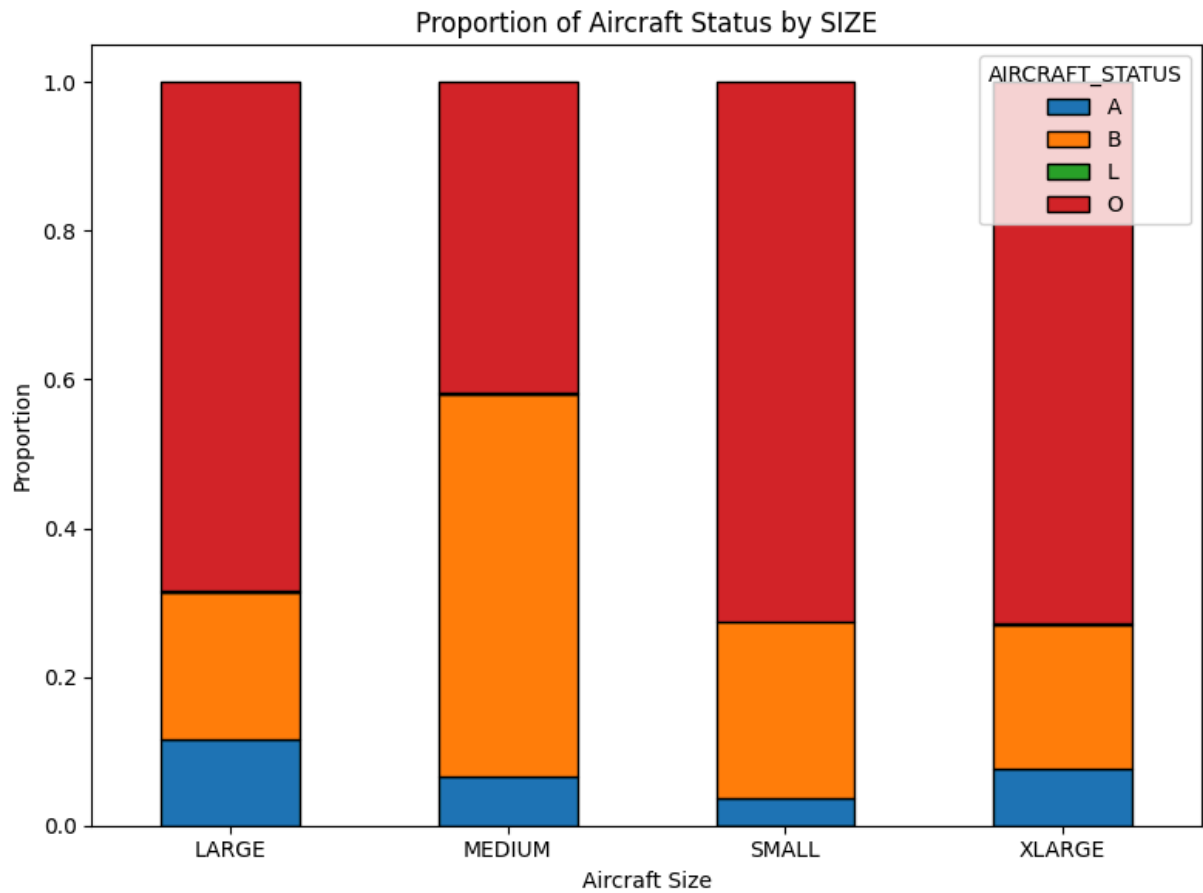
```
In [36]: import matplotlib.pyplot as plt

# Grouping by size and aircraft_status
status_summary = (
    inventory_cleaned.groupby(['SIZE', 'AIRCRAFT_STATUS'])
    .size()
    .unstack(fill_value=0)
)

# Converting proportions
status_proportions = status_summary.div(status_summary.sum(axis=1), axis=0)

# Plotting
status_proportions.plot(kind='bar', stacked=True, figsize=(8, 6), edgecolor=

plt.title('Proportion of Aircraft Status by SIZE')
plt.xlabel('Aircraft Size')
plt.ylabel('Proportion')
plt.legend(title='AIRCRAFT_STATUS')
plt.xticks(rotation=0)
plt.tight_layout()
plt.show()
print(status_proportions)
```



AIRCRAFT_STATUS	A	B	L	O
SIZE				
LARGE	0.115702	0.197468	0.001746	0.685084
MEDIUM	0.065880	0.514826	0.001096	0.418198
SMALL	0.037909	0.235739	0.000000	0.726352
XLARGE	0.076263	0.193708	0.001732	0.728296

From these plots we can first see that the majority of aircraft across all size groups fall under 'Y' operating status and the 'O' and 'B' aircraft status categories. Small, large, and extra-large aircrafts all had only about 3% of their aircrafts not operational, while medium aircrafts were slightly higher at 5%. This is interesting because it correlates to what we find in the aircraft status plot as well. Medium aircrafts had the lowest 'O' status and highest 'B' status compared to the other sized aircrafts. This plot also seems to show a correlation between aircraft size and 'A' status, for as the planes get larger, the 'A' status frequency increases.